

```
pip install gradio
```

```

46.5/46.5 kB 3.0 MB/s eta 0:00:00
Downloading mdit_py_plugins-0.3.0-py3-none-any.whl (43 kB)
43.7/43.7 kB 3.4 MB/s eta 0:00:00
Downloading mdit_py_plugins-0.2.8-py3-none-any.whl (41 kB)
41.0/41.0 kB 2.6 MB/s eta 0:00:00
Downloading mdit_py_plugins-0.2.7-py3-none-any.whl (41 kB)
41.0/41.0 kB 3.1 MB/s eta 0:00:00
Downloading mdit_py_plugins-0.2.6-py3-none-any.whl (39 kB)
Downloading mdit_py_plugins-0.2.5-py3-none-any.whl (39 kB)
INFO: pip is looking at multiple versions of mdit-py-plugins to determine which version is compatible with other requirements. This could take a while.
Downloading mdit_py_plugins-0.2.4-py3-none-any.whl (39 kB)
Downloading mdit_py_plugins-0.2.3-py3-none-any.whl (39 kB)
Downloading mdit_py_plugins-0.2.2-py3-none-any.whl (39 kB)
Downloading mdit_py_plugins-0.2.1-py3-none-any.whl (38 kB)
Downloading mdit_py_plugins-0.2.0-py3-none-any.whl (38 kB)
INFO: This is taking longer than usual. You might need to provide the dependency resolver with stricter constraints to reduce runtime. See https://pip.pypa.io/warnings/backtracking
Downloading mdit_py_plugins-0.1.0-py3-none-any.whl (37 kB)
Collecting markdown-it-py[linkify]>=2.0.0 (from gradio)
Downloading markdown_it_py-3.0.0-py3-none-any.whl (87 kB)
87.5/87.5 kB 7.1 MB/s eta 0:00:00
Downloading markdown_it_py-2.2.0-py3-none-any.whl (84 kB)
84.5/84.5 kB 7.3 MB/s eta 0:00:00
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/dist-packages (from pandas<3.0,>=1.0->gradio) (2022.7.1)
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.10/dist-packages (from requests~=2.0->gradio) (3.4)
Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.10/dist-packages (from requests~=2.0->gradio) (1.26.16)
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.10/dist-packages (from requests~=2.0->gradio) (2023.7.22)
Requirement already satisfied: click>=7.0 in /usr/local/lib/python3.10/dist-packages (from uvicorn>=0.14.0->gradio) (8.1.6)
Collecting h11>=0.8 (from uvicorn>=0.14.0->gradio)
Downloading h11-0.14.0-py3-none-any.whl (58 kB)
58.3/58.3 kB 4.3 MB/s eta 0:00:00
Collecting starlette<0.28.0,>=0.27.0 (from fastapi->gradio)
Downloading starlette-0.27.0-py3-none-any.whl (66 kB)
67.0/67.0 kB 4.0 MB/s eta 0:00:00
Collecting httpcore<0.18.0,>=0.15.0 (from httpx->gradio)
Downloading httpcore-0.17.3-py3-none-any.whl (74 kB)
74.5/74.5 kB 4.0 MB/s eta 0:00:00
Requirement already satisfied: sniffio in /usr/local/lib/python3.10/dist-packages (from httpx->gradio) (1.3.0)
Requirement already satisfied: anyio<5.0,>=3.0 in /usr/local/lib/python3.10/dist-packages (from httpcore<0.18.0,>=0.15.0->httpx->gradio) (3.7.1)
Requirement already satisfied: pyrsistent!=0.17.0,!=0.17.1,!=0.17.2,>=0.14.0 in /usr/local/lib/python3.10/dist-packages (from jsonschema>=3.0->altair<6.0,>=4.2.0->gradio) (0.19.3)
Requirement already satisfied: uc-micro-py in /usr/local/lib/python3.10/dist-packages (from linkify-it-py<3,>=1->markdown-it-py[linkify]>=2.0.0->gradio) (1.0.2)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.10/dist-packages (from python-dateutil>=2.7->matplotlib~=3.0->gradio) (1.16.0)
Requirement already satisfied: exceptiongroup in /usr/local/lib/python3.10/dist-packages (from anyio<5.0,>=3.0->httpcore<0.18.0,>=0.15.0->httpx->gradio) (1.1.2)
Building wheels for collected packages: ffmpeg
  Building wheel for ffmpeg (setup.py) ... done
  Created wheel for ffmpeg: filename=ffmpeg-0.3.1-py3-none-any.whl size=5579 sha256=1571b9899afc5bc5136ee65fc6d47497e8335946927c0d9d42c6cf461ac9e5aa
  Stored in directory: /root/.cache/pip/wheels/01/a6/d1/1c0828c304a4283b2c1639a09ad86f83d7c487ef34c6b4a1bf
Successfully built ffmpeg
Installing collected packages: pydub, ffmpeg, websockets, semantic-version, python-multipart, orjson, markdown-it-py, h11, aiofiles, uvicorn, starlette, mdit-py-plugins, huggingface
Attempting uninstall: markdown-it-py
  Found existing installation: markdown-it-py 3.0.0
  Uninstalling markdown-it-py-3.0.0:
    Successfully uninstalled markdown-it-py-3.0.0
Attempting uninstall: mdit-py-plugins
  Found existing installation: mdit-py-plugins 0.4.0
  Uninstalling mdit-py-plugins-0.4.0:
    Successfully uninstalled mdit-py-plugins-0.4.0
Successfully installed aiofiles-23.1.0 fastapi-0.101.0 ffmpeg-0.3.1 gradio-3.39.0 gradio-client-0.3.0 h11-0.14.0 httpcore-0.17.3 httpx-0.24.1 huggingface-hub-0.16.4 markdown-it-py-2

```

```
!pip install openai
```

Collecting openai

Downloading openai-0.27.8-py3-none-any.whl (73 kB)

73.6/73.6 kB 2.0 MB/s eta 0:00:00

Requirement already satisfied: requests>=2.20 in /usr/local/lib/python3.10/dist-packages (from openai) (2.31.0)

Requirement already satisfied: tqdm in /usr/local/lib/python3.10/dist-packages (from openai) (4.65.0)

Requirement already satisfied: aiohttp in /usr/local/lib/python3.10/dist-packages (from openai) (3.8.5)

Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.10/dist-packages (from requests>=2.20->openai) (3.2.0)

Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.10/dist-packages (from requests>=2.20->openai) (3.4)

Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.10/dist-packages (from requests>=2.20->openai) (1.26.16)

Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.10/dist-packages (from requests>=2.20->openai) (2023.7.22)

Requirement already satisfied: attrs>=17.3.0 in /usr/local/lib/python3.10/dist-packages (from aiohttp->openai) (23.1.0)

Requirement already satisfied: multidict<7.0,>=4.5 in /usr/local/lib/python3.10/dist-packages (from aiohttp->openai) (6.0.4)

Requirement already satisfied: async-timeout<5.0,>=4.0.0a3 in /usr/local/lib/python3.10/dist-packages (from aiohttp->openai) (4.0.2)

Requirement already satisfied: yarl<2.0,>=1.0 in /usr/local/lib/python3.10/dist-packages (from aiohttp->openai) (1.9.2)

Requirement already satisfied: frozenlist>=1.1.1 in /usr/local/lib/python3.10/dist-packages (from aiohttp->openai) (1.4.0)

Requirement already satisfied: aiosignal>=1.1.2 in /usr/local/lib/python3.10/dist-packages (from aiohttp->openai) (1.3.1)

Installing collected packages: openai

Successfully installed openai-0.27.8

```
import gradio as gr
```

```
import openai
```

```
import os
```

```
import warnings
```

```
warnings.filterwarnings("ignore")
```

```
#Add the access key generated from your OpenAI account
```

```
openai.api_key = "sk-JDqGzBGpGErhFHfjHwM7T3B1bkFJRwWHz9JLqsi8XSNPkr6N"
```

```
def chatgpt_api(input_text):
```

```
    messages = [
```

```
        {"role": "system", "content": "You are a helpful assistant."}]
```

```
    if input_text:
```

```
        messages.append(
```

```
            {"role": "user", "content": 'Summarize this text "{}" into a short and concise Dall-e2 prompt'.format(input_text)},
```

```
        )
```

```
        chat_completion = openai.ChatCompletion.create(
```

```
            model="gpt-3.5-turbo", messages=messages
```

```
        )
```

```
    reply = chat_completion.choices[0].message.content
```

```
    return reply
```

```
def dall_e_api(dalle_prompt):
```

```
    dalle_response = openai.Image.create(
```

```
        prompt = dalle_prompt,
```

```
        size="512x512"
```

```
    )
```

```
    image_url = dalle_response['data'][0]['url']
```

```
    return image_url
```

```
def whisper_transcribe(audio):
```

```
    os.rename(audio, audio + '.wav')
```

```
audio_file = open(audio + '.wav', "rb")
transcript = openai.Audio.transcribe("whisper-1", audio_file)
dalle_prompt = chatgpt_api(transcript["text"])
image_url = dall_e_api(dalle_prompt)
return transcript["text"], image_url
```

```
output_1 = gr.Textbox(label="Speech to Text")
output_2 = gr.Image(label="DALL-E Image")
```

```
speech_interface = gr.Interface(fn = whisper_transcribe,
                                inputs = gr.Audio(source="microphone", type="filepath"),
                                outputs = [output_1, output_2],
                                title = "Generate Images using Voice")
```

```
speech_interface.launch(debug=True)
```

... Colab notebook detected. This cell will run indefinitely so that you can see errors and logs. To turn off, set debug=False. Note: opening Chrome Inspector may crash demo inside Colab notebooks.

To create a public link, set `share=True` in `launch()`.
Running on <https://localhost:7860/>

Generate Images using Voice

🎵 audio

0:00


Clear

Submit

Speech to Text

Generate a Portrait of an Old Woman in Dress, Photorealistic and HD

📎 DALL-E Image



Download

[Colab paid products](#) - [Cancel contracts here](#)

Executing (4m 46s) <cell line: 6> > launch() > block_thread()

